



A jumpstart into machine-learning

Dahvyd Wing

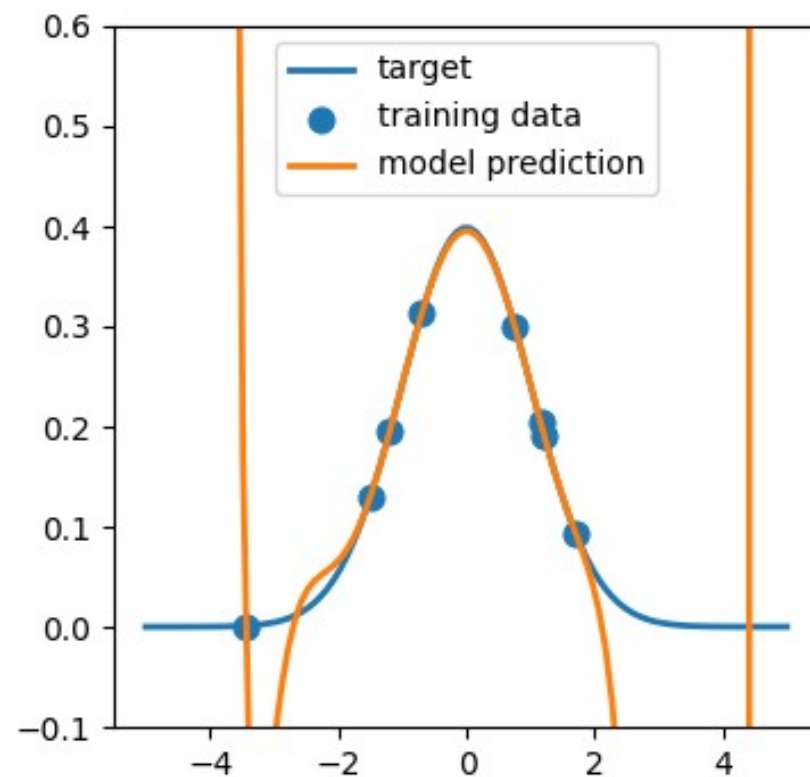


Regression for computational chemistry

Most applications: lots of noisy data



Our case: few data points with almost no noise



Anatomy of regression


1. Data:

2. Model:

are trainable parameters

3. Cost function:
mean squared error (MSE)

4. Find using gradient descent




Pytorch example lj_1_overfit.py

1. Open anaconda prompt
2. conda activate ml_tutorial
3. Go to ml_tutorial folder
4. spyder &
5. In spyder open lj_1_overfit.py

Anatomy of regression

1. Data:

- Instance/object of a customized dataset class
- Implement 3 functions: `__init__`, `__len__`, and `__getitem__`
- dataloader pulls random batches of data from the dataset



Anatomy of regression


1. Data:

2. Model:

- Instance/object of a customized nn.module class
- Implement 2 functions: `__init__` and `__forward__`
- Descriptor:

One hot encoding:

Continuous data: $x = 3.1$ (→ 1 0.5 0.3 0.01 0 0)



Anatomy of regression


1. Data:

2. Model:

- Instance/object of a customized nn.module class
- Implement 2 functions: `__init__` and `__forward__`
- Descriptor
- Neural network:

$$y_{pred} = \mathbf{w}_3 \cdot \mathbf{y}_3 + b_3$$

is the nonlinear activation function: GELU



Use a continuously differentiable activation function

Anatomy of regression

1. Data:

2. Model:

3. Cost function:

- Mean squared error

Anatomy of regression

1. Data:
2. Model:
3. Cost function:
4. Find

Batch 1

, ,

$$\boldsymbol{\theta}' = \boldsymbol{\theta} - \nabla_{\boldsymbol{\theta}} C(\mathbf{y}, f(\mathbf{X}; \boldsymbol{\theta}))$$

Batch 2

, ,

$$\boldsymbol{\theta}' = \boldsymbol{\theta} - \nabla_{\boldsymbol{\theta}} C(\mathbf{y}, f(\mathbf{X}; \boldsymbol{\theta}))$$

Batch 3

, ,



$$\boldsymbol{\theta}' = \boldsymbol{\theta} - \nabla_{\boldsymbol{\theta}} C(\mathbf{y}, f(\mathbf{X}; \boldsymbol{\theta}))$$

...

1 Epoch

Overfitting



- NNs often have many more parameters than samples in the training data
- Run `lj_1_overfit.py` several times
 - 6,601 trainable parameters, 10 data points



- Each model perfectly fits the training points, but doesn't do a great job in the interpolation region
You measure a model by testing on data it has never seen

Overfitting


- NNs often have many more parameters than samples in the training data
- Run `lj_1_overfit.py` several times
 - 6,601 trainable parameters, 10 data points



- Each model perfectly fits the training points, but doesn't do a great job in the interpolation region
- The models do terrible in the extrapolation regime

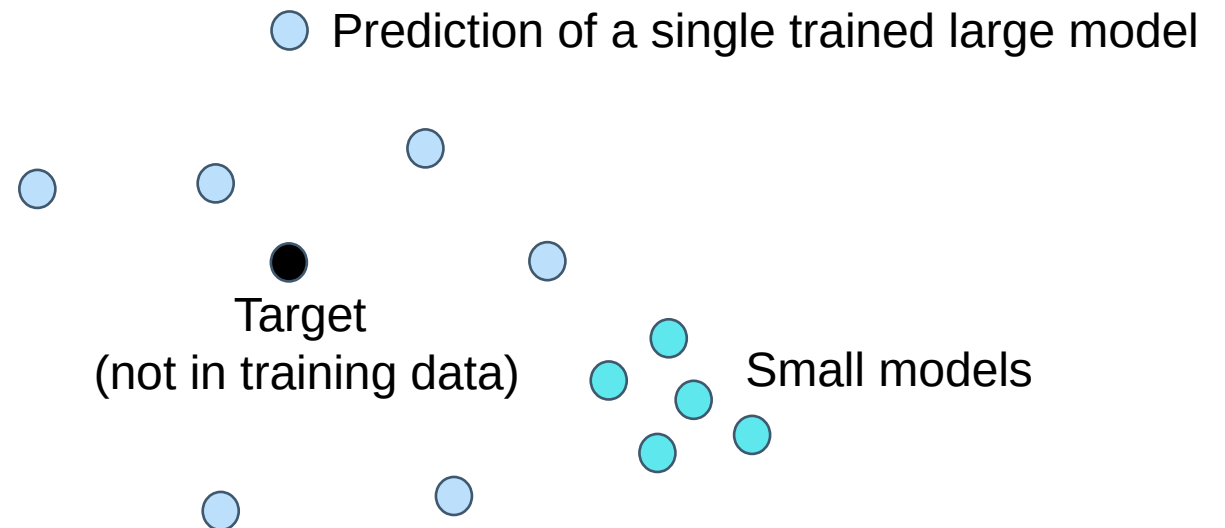

Validation and test sets

- Separate your data into a training set, a validation set, and a test set
- Validation set used to measure overfitting and tune hyperparameters
- Test set is only used for the final model to get a final estimate of how accurate the model really is
- Run `lj_2_overfit_with_validation.py`
 - The main change in the code is lines 56-58
- To get the best performance model stop when there is a steady increase in validation loss and decrease in training loss (early stopping)




Learning curve

- With enough data the validation loss and training loss should converge
- Variance: a models trained with different, but equal number of points yield different results
 - The more parameters, the more variance in the model
- Large high variance models overfit



Learning curve


- With enough data the validation loss and training loss should converge
- Variance: a models trained with different, but equal number of points yield different results
 - The more parameters, the more variance in the model
- Large high variance models overfit



- Enough data removes variance
- We are always working in the low data regime

Bias variance tradeoff

- There is an optimum size of your model for a given amount of data.



- L2 regularization lowers variance.
 - Prevents overfitting
 - Allows you to use larger models

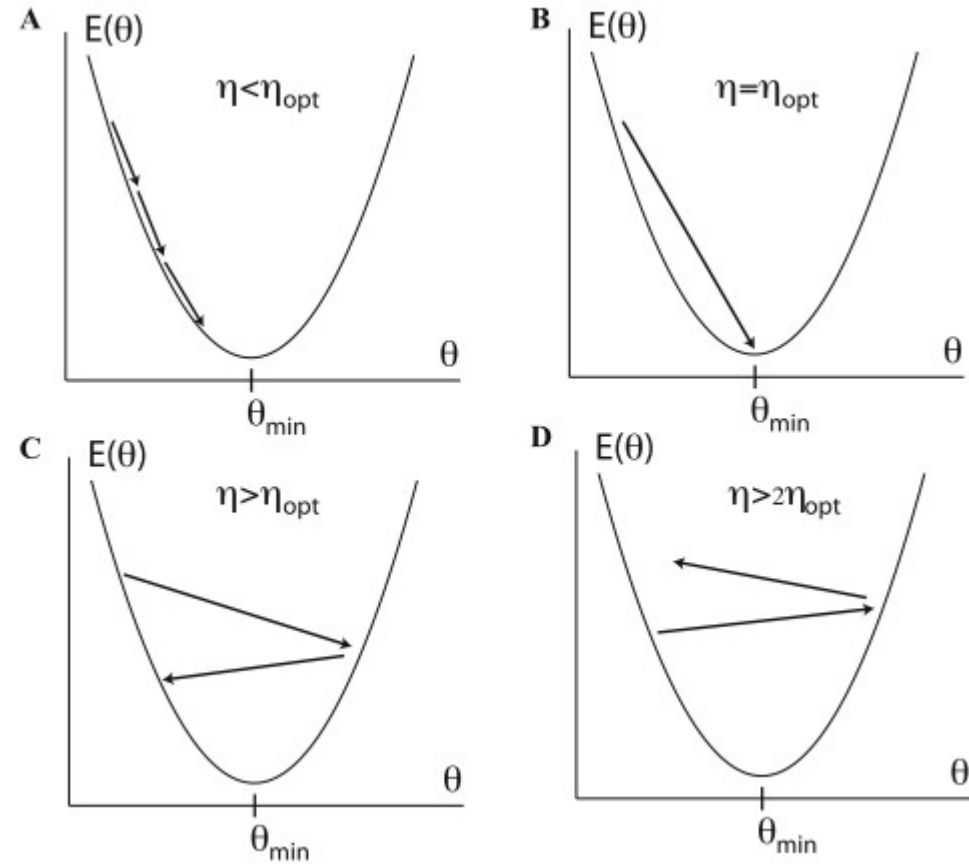
Use the right amount of regularization

Find : Gradient descent

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \mathbf{v}_t$$

- Momentum algorithms
 - Build up speed in shallow directions

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \mathbf{v}_t$$




Momentum

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \mathbf{v}_t$$

- Momentum algorithms
 - Build up speed in shallow directions
 - Can get out of local minima

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \mathbf{v}_t$$




Use stochasticity to get out of local minima

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \mathbf{v}_t$$

Only compute on a subset of \mathbf{x} and \mathbf{y}

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \mathbf{v}_t$$



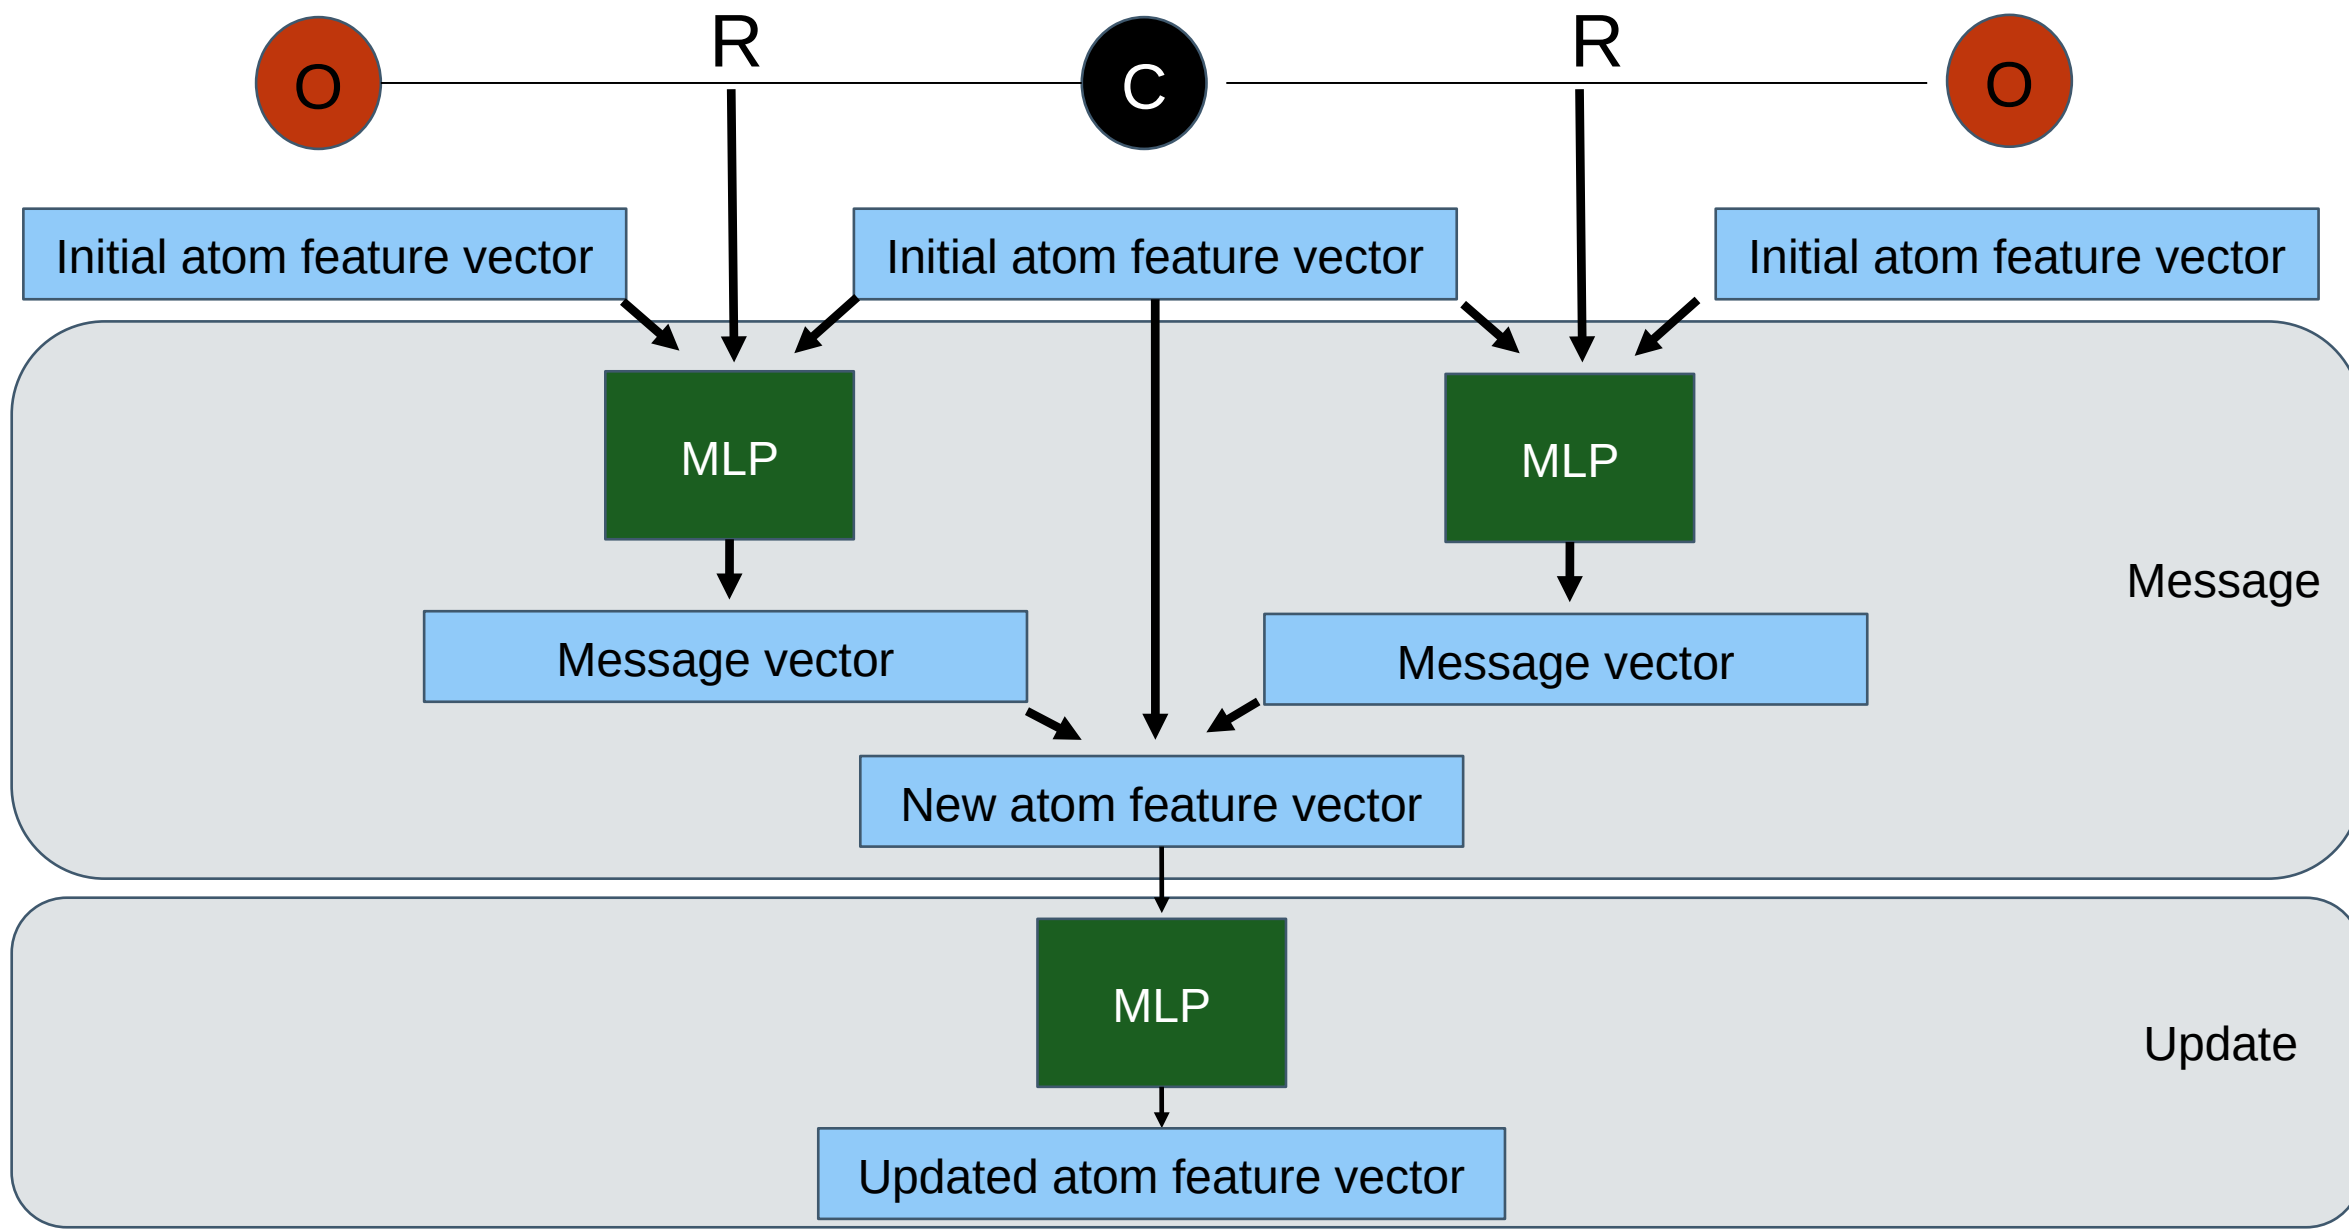
A model to play with

- Run `lj_3_hyperparameters.py`
- Change hyperparameters at the top of the script and see how the training progression changes
- Tensorboard to plot training progression
 - In `ml_tutorial` folder run: `tensorboard --logdir=runs --reload_multifile True`
 - Go to <http://localhost:6006/> in your browser

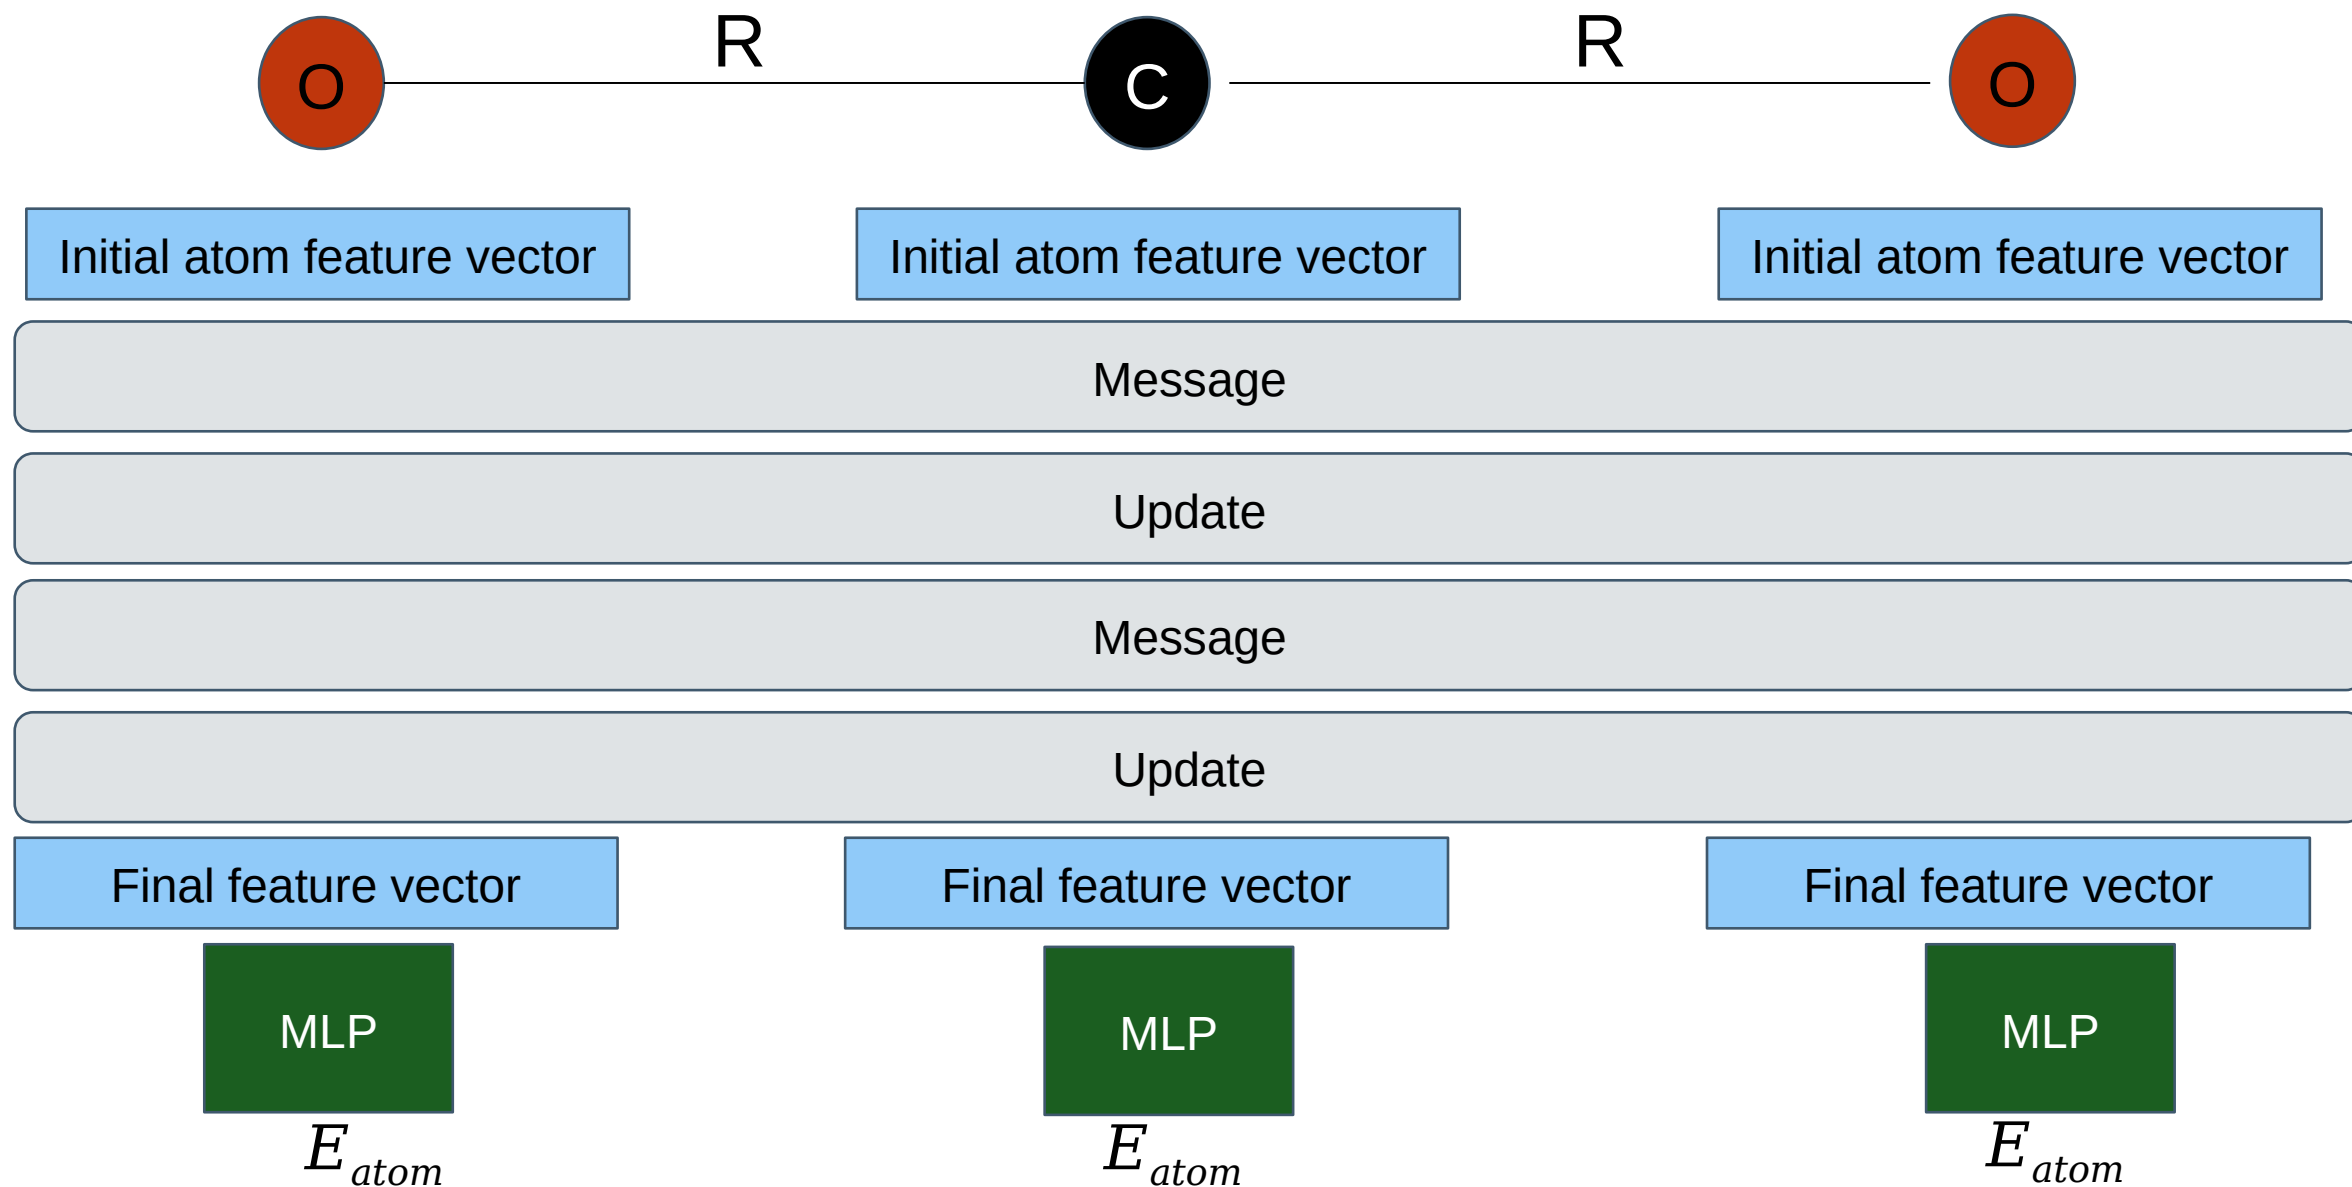
Best practices when developing an NN

- Try to memorize a few data points first
- Try also using fake simple data
- Check your descriptors and targets to make sure you are feeding the NN what you think you are
- Hyperparameter tuning on even just a few epochs to screen out unpromising parameter values
- Use grid search/random search of hyperparameters

Architecture of message passing neural networks



Architecture of message passing neural networks



Getting to molecular dynamics in 1 minute

- In `schmet_tutorial` folder run `script2_train_schmet.py`
 - Trains a schmet model on 1000 data points of a ethanol DFT(PBE) MD trajectory for 5 epochs
- run `script3_run_md.py`
 - Creates an atomic structure environment (ASE) force field calculator using the schmet model. Uses ASE to run an MD trajectory
 - Note: that you can only run this script once, it will throw an error if you run it a second time
- in anaconda prompt run: `ase gui`
- file -> open -> `ase_calcs` (in left pane) -> `simulation.traj`
- tools -> movies -> play
- Note that the state-of-the-art ML-FF's are equivariant (PAINN, NequIP, etc.). Schmet is an invariant model that's cheap to train and run.


Acknowledgements



Prof. Alexandre Tkatchenko



Artem Kokorin



Funded by the
European Union



Thank You!

